

HBASE**Exercise 1:****COMMAND:**

```
create 'csp554Tbl', 'cf1', 'cf2'

describe 'csp554Tbl'
```

Output:

```
COLUMN FAMILIES DESCRIPTION
{NAME => 'cf1', BLOOMFILTER => 'ROW', IN_MEMORY => 'false', VERSIONS => '1', KEEP_DELETED_CELLS => 'FALSE',
E => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}

{NAME => 'cf2', BLOOMFILTER => 'ROW', IN_MEMORY => 'false', VERSIONS => '1', KEEP_DELETED_CELLS => 'FALSE',
E => 'true', BLOCKSIZE => '65536', REPLICATION_SCOPE => '0'}
```

Exercise 2:**COMMAND:**

```
scan 'csp554Tbl'
```

OUTPUT:

```
hbase:011:0> scan 'csp554Tbl'
ROW                                COLUMN+CELL
Row1                                column=cf1:name, timestamp=2023-11-30T23:35:38.573, value=Sam
Row1                                column=cf2:job, timestamp=2023-11-30T23:35:56.954, value=Pilot
Row1                                column=cf2:level, timestamp=2023-11-30T23:36:11.747, value=LZ3
Row2                                column=cf1:name, timestamp=2023-11-30T23:35:48.915, value=Ahmed
Row2                                column=cf2:job, timestamp=2023-11-30T23:36:04.766, value=Doctor
Row2                                column=cf2:level, timestamp=2023-11-30T23:36:18.995, value=AR7
2 row(s)
Took 0.1280 seconds
hbase:012:0> |
```

Exercise 3:**COMMAND:**

```
get 'csp554Tbl', 'Row1', 'cf2:level'
```

OUTPUT:

```
hbase:012:0> get 'csp554Tbl', 'Row1', 'cf2:level'
COLUMN                                CELL
cf2:level                             timestamp=2023-11-30T23:36:11.747, value=LZ3
1 row(s)
Took 0.0437 seconds
hbase:013:0> |
```

Exercise 4:**COMMAND:**

```
get 'csp554Tbl', 'Row2', 'cf1:name'
```

OUTPUT:

```
hbase:013:0> get 'csp554Tbl', 'Row2', 'cf1:name'
COLUMN                                CELL
cf1:name                              timestamp=2023-11-30T23:35:48.915, value=Ahmed
1 row(s)
Took 0.0178 seconds
hbase:014:0>
```

Exercise 5:**COMMAND:**

scan 'csp554Tbl', {LIMIT => 2}

OUTPUT:

```
hbase:014:0> scan 'csp554Tbl', {LIMIT => 2}
ROW                                     COLUMN+CELL
Row1                                  column=cf1:name, timestamp=2023-11-30T23:35:38.573, value=Sam
Row1                                  column=cf2:job, timestamp=2023-11-30T23:35:56.954, value=Pilot
Row1                                  column=cf2:level, timestamp=2023-11-30T23:36:11.747, value=LZ3
Row2                                  column=cf1:name, timestamp=2023-11-30T23:35:48.915, value=Ahmed
Row2                                  column=cf2:job, timestamp=2023-11-30T23:36:04.766, value=Doctor
Row2                                  column=cf2:level, timestamp=2023-11-30T23:36:18.995, value=AR7
2 row(s)
Took 0.0341 seconds
hbase:015:0>
```

Cassandra**Exercise 1:**

ANS : The article proposes a novel, query-driven data modeling methodology tailored for the Apache Cassandra NoSQL database. Traditional relational database modeling focuses primarily on conceptual data modeling and normalization. In contrast, this Cassandra methodology puts equal emphasis on understanding the application queries that will be run against the database. It has four main components:

- ✓ Conceptual data modeling to understand the data
- ✓ Application workflow modeling to understand the access patterns/queries
- ✓ Logical data modeling driven by the queries to design efficient Cassandra table schemas
- ✓ Physical data modeling to optimize data types, keys, partitions etc.

The methodology relies heavily on principles like data nesting within table partitions and intentional data duplication across tables to optimize for fast writes and reads in Cassandra. The authors introduce new modeling principles, mapping rules, patterns and visualization diagrams to guide Cassandra schema design. They also present a tool, KDM, that automates major parts of the methodology.

Comments:

This methodology seems like a rigorous, systematic approach tailored to leverage Cassandra's architecture. Mapping application access patterns directly into the logical schema design is smart. The conceptual diagrams provide useful visualization. Automation through the KDM tool would be extremely helpful for developers. My main question would be how adaptable this is to schema changes over time as application queries evolve. But overall this appears to be a solid modeling foundation for Cassandra databases.

Exercise 2:

COMMAND:

DESCRIBE TABLE MUSIC;

OUTPUT:

```
cqlsh:a20528281> DESCRIBE TABLE Music;

CREATE TABLE a20528281.music (
  artistname text,
  albumname text,
  cost int,
  numbersold int,
  PRIMARY KEY (artistname, albumname)
) WITH CLUSTERING ORDER BY (albumname ASC)
  AND bloom_filter_fp_chance = 0.01
  AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
  AND comment = ''
  AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
  AND compression = {'chunk_length_in_kb': '64', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
  AND crc_check_chance = 1.0
  AND dclocal_read_repair_chance = 0.1
  AND default_time_to_live = 0
  AND gc_grace_seconds = 864000
  AND max_index_interval = 2048
  AND memtable_flush_period_in_ms = 0
  AND min_index_interval = 128
  AND read_repair_chance = 0.0
  AND speculative_retry = '99PERCENTILE';

cqlsh:a20528281> |
```

Exercise 3:

CODE :

```
INSERT INTO Music (artistName, albumName, numberSold, cost)
VALUES
```

```
('Mozart', 'Greatest Hits', 100000, 10),
('Taylor Swift', 'Fearless', 2300000, 15),
('Black Sabbath', 'Paranoid', 534000, 12),
('Katy Perry', 'Prism', 800000, 16),
('Katy Perry', 'Teenage Dream', 750000, 14);
```

COMMAND:

```
SELECT * FROM Music;
```

OUTPUT:

```
cqlsh:a20528281> SELECT * FROM Music;

artistname | albumname | cost | numbersold
-----+-----+-----+-----
Mozart      | Greatest Hits | 10   | 100000
Black Sabbath | Paranoid | 12   | 534000
Taylor Swift | Fearless | 15   | 2300000
Katy Perry  | Prism | 16   | 800000
Katy Perry  | Teenage Dream | 14   | 750000

(5 rows)
cqlsh:a20528281>
```

Exercise 4 :**COMMAND:**

```
source 'ex4.cql';
```

Output:

```
cqlsh:a20528281> source 'ex4.cql';

artistname | albumname | cost | numbersold
-----+-----+-----+-----
Katy Perry | Prism | 16   | 800000
Katy Perry | Teenage Dream | 14   | 750000

(2 rows)
cqlsh:a20528281>
```

Exercise 5:**COMMAND:**

```
Source 'ex5.cql';
```

Output:

```
(2 rows)
cqlsh:a20528281> source 'ex5.cql';

  artistname | albumname | cost | numbersold
-----+-----+-----+-----
  Taylor Swift | Fearless | 15 | 2300000
    Katy Perry | Prism | 16 | 800000
    Katy Perry | Teenage Dream | 14 | 750000

(3 rows)
cqlsh:a20528281>
```

MONGO DB**Exercise 1:****COMMAND:**

```
db.unicorns.find({ weight: { $lt: 500 } });
```

Output:

```
> db.unicorns.find({ weight: { $lt: 500 } });
{ "_id" : ObjectId("656a41166feda63e9ff3de87"), "name" : "Aurora", "dob" : ISODate("1991-01-24T13:00:00Z"), "loves" : [ "carrot", "grape" ], "weight" : 450, "gender" : "f", "vampires" : 43 }
{ "_id" : ObjectId("656a41176feda63e9ff3de8d"), "name" : "Raleigh", "dob" : ISODate("2005-05-03T00:57:00Z"), "loves" : [ "apple", "sugar" ], "weight" : 421, "gender" : "m", "vampires" : 2 }
>
```

Exercise 2:**Command:**

```
db.unicorns.find({ loves: "apple" });
```

Output:

```
db.unicorns.find({ loves: "apple" });
{ "_id" : ObjectId("656a41166feda63e9ff3de89"), "name" : "Roosnoodes", "dob" : ISODate("1979-08-18T18:44:00Z"), "loves" : [ "apple" ], "weight" : 575, "gender" : "m", "vampires" : 99 }
{ "_id" : ObjectId("656a41166feda63e9ff3de8a"), "name" : "Solnara", "dob" : ISODate("1985-07-04T02:01:00Z"), "loves" : [ "apple", "carrot", "chocolate" ], "weight" : 550, "gender" : "f", "vampires" : 80 }
{ "_id" : ObjectId("656a41176feda63e9ff3de8d"), "name" : "Raleigh", "dob" : ISODate("2005-05-03T00:57:00Z"), "loves" : [ "apple", "sugar" ], "weight" : 421, "gender" : "m", "vampires" : 2 }
{ "_id" : ObjectId("656a41176feda63e9ff3de8e"), "name" : "Leia", "dob" : ISODate("2001-10-08T14:53:00Z"), "loves" : [ "apple", "watermelon" ], "weight" : 601, "gender" : "f", "vampires" : 33 }
{ "_id" : ObjectId("656a41176feda63e9ff3de8f"), "name" : "Pilot", "dob" : ISODate("1997-03-01T05:03:00Z"), "loves" : [ "apple", "watermelon" ], "weight" : 650, "gender" : "m", "vampires" : 54 }
```

Exercise 3:**Command:**

```
db.unicorns.insert({
  name: "Malini",
  dob: new Date("2008-03-11T00:00:00Z"),
```

```
    loves: ["pears", "grapes", "carrots"],
    weight: 470,
    gender: "F",
    vampires: 24,
    horns: 1
  });

  db.unicorns.find({ name: "Malini" });
```

Output:

```
writeResult({ "inserted" : 1 })
> db.unicorns.find({ name: "Malini" });
{ "_id" : ObjectId("656a4a26feda63e9ff3de93"), "name" : "Malini", "dob" : ISODate("2008-03-11T00:00:00Z"), "loves" : [ "pears", "grapes" ], "weight" : 450, "gender" : "F", "vampires" : 23, "horns" : 1 }
```

Exercise 4:

Command:

```
    db.unicorns.update(
    { name: "Malini" },
    { $push: { loves: "grapes" } }
  );

  db.unicorns.find({ name: "Malini" }, { loves: 1, _id: 0 });
```

Output:

```
writeResult({ "matched" : 1, "upserted" : 0, "modified" : 1 })
> db.unicorns.find({ name: "Malini" }, { loves: 1, _id: 0 });
{ "loves" : [ "pears", "grapes", "apricots" ] }
> .....
```

Exercise 5:

Command:

```
db.unicorns.deleteMany({ weight: { $gt: 600 } });
```

Output:

```
> db.unicorns.deleteMany({ weight: { $gt: 600 } });
{ "acknowledged" : true, "deletedCount" : 6 }
> .....
```